

AI Disclosure Statement — HoneyMoney

Prepared by: Team JUST50 · Last updated: 23 August 2026

HoneyMoney is an AI-driven financial wellness application for Malaysian households, built on the philosophy of “Happy Wife, Happy Life”. It gives a couple funding transparency where money is shared and spending autonomy where it is personal, without the tracking fatigue that makes conventional budgeting apps fail. Because it handles people's money, transparency about *how* and *where* we use artificial intelligence is a condition of trust. This statement discloses the AI in our product, the AI we used to build it, and the limits we have put on both.

Where a claim is an estimate rather than a measured fact, we say so, and we have removed claims this document previously made that the shipped product does not support. A capability that survives in a disclosure after it has left the code is the same defect as an undisclosed one.

1. The problem, and where AI actually helps

Malaysian households run money across a fragmented payment landscape — Touch 'n Go eWallet, Maybank MAE, GrabPay, ShopeePay, cards and cash — and nothing adds them up. Conventional apps answer this with manual data entry and surveillance, which people abandon, and they store the result in flat tables that cannot warn anyone before a goal is missed. HoneyMoney models the household as a **living knowledge graph** of how income flows, allocated across three buckets: **Must-Paid** (shared essentials, visible to both partners), **Savings** (a percentage moved before it can be spent), and **Privacy** (personal spending where tracking stops).

AI does exactly two jobs in that system. It is not the system.

Job	What the model does	What it is never allowed to do
Reading a receipt	Extracts merchant, amount, date and a suggested category from a forwarded e-wallet screenshot or a photographed receipt.	Write to the ledger. Every parsed transaction is shown back for explicit human confirmation first.
Explaining the numbers	Phrases figures that have already been computed into plain-language guidance, as our persona “Honey”.	Calculate any figure a user sees, or recommend a financial product.

2. The AI explains the money. Exact code computes it.

Honey uses AI for what language models are genuinely excellent at: making sense of messy real-world input — a photographed receipt, a forwarded screenshot, a sentence typed one-handed — and explaining a household's position in plain English or Malay. That is real work, and it is why the product feels like talking to someone rather than filling in a form.

The figures in that explanation come from a deterministic function — the same one behind the H-Score screen — so a household gets both halves at once: **language that sounds human, and arithmetic that is exact and reproducible**. A final check confirms the wording matches the computed figures before anyone sees it, so the two can never drift apart. Advice a family can act on has to be both understandable and correct; this is how HoneyMoney delivers both rather than trading one for the other.

Ask Honey runs in three stages: **parse** → **compute** → **narrate**.

- **Stage 2 owns every number the user ever sees.** It is a pure function: no database, no network, no model, no clock. “What would this purchase do to us?” is answered by applying the purchase and calling `computeHScore` — the identical function behind the H-Score screen, not a parallel approximation of it. Two engines that disagree would tell one household two different truths with no way to tell which was real.
- **Stage 3 turns those numbers into language a family actually uses** — and every figure it writes is matched back against the ones stage 2 published before the answer is shown. The guarantee this buys is a strong one and easy to state: **every number a household reads has been computed, never composed**. Where the check

does not pass, Honey answers from a deterministic template, so the household still gets an answer and it is still exact.

- **Honey's scope is defined in code, so it stays a coach.** The questions it answers are an explicit set built around understanding your own money. "Which unit trust should I buy?" is routed to a licensed directory, because product recommendation is regulated advice and belongs with people who are licensed to give it. That boundary is structural rather than a prompt instruction, which is what makes it dependable.
- **Honey asks rather than assumes.** An affordability question with no price gets a question back instead of a guessed figure — a guessed price would quietly turn a budgeting tool into a product recommender, and a price looked up online is someone else's TV, not yours. Asking takes one extra tap and keeps the answer about your household.

The 3-Bucket allocation is rules-based, not AI-decided. Money moves between Must-Paid, Savings and Privacy by user-configured rules. The AI helps people *understand* their money; it does not autonomously move or lock funds, and it does not execute transactions.

3. Which models, and what they receive

HoneyMoney uses **switchable AI providers**, selected per deployment, with a per-call token and cost ledger for governance. Three are implemented and can be probed live at `/api/ai/check` :

Provider	Role	Where it runs
Google Gemini Flash	Multimodal — the only provider that also does image OCR	Cloud
Groq	Text insight generation	Cloud
Ollama (local)	Text insight generation	On-device — nothing leaves the machine

Switchability is a resilience and cost property, not a marketing one: no vendor lock-in, and a free-tier outage at one provider does not take the product down. Adding further providers is a roadmap item, not a shipped capability.

A household may supply its own key, so the AI works without depending on whoever runs the server. Because that key is a live billable credential stored beside a family's ledger, it is **encrypted (AES-256-GCM) by the application before the database sees it**, is never sent back to a browser, and is validated against the provider at the moment it is saved. If no encryption key is configured on the server, the app *refuses to store* the credential rather than falling back to storing it in the clear.

Receipt reading has a zero-token path. The default in-app receipt scan runs **on-device OCR** (Tesseract in the browser) and never contacts a model at all. The cloud model is used for the Telegram forwarding path and where on-device extraction is not good enough.

What the model is sent. For extraction: the forwarded image, returning structured text. For insight: *derived, minimised* figures — aggregates and category summaries needed to answer the question — never a raw dump of a household's ledger. We do not send images or data to any model for advertising, model training, or any purpose beyond serving the user's request.

4. Privacy, and a measured privacy fix

- **Marital-safe by design.** The Privacy bucket is private. Honey is constrained never to expose one partner's private-wallet detail to the other — and that constraint is enforced by the same query filter the on-screen record list uses, not by a prompt. **This was tested to discriminate rather than pass vacuously:** with the filter, partner B is told RM4,634; without it, RM991,158. A check that cannot fail against the old behaviour proves nothing.
- **Household-level figures stay household-level on purpose.** H-Score, income and bucket allocations are visible to both partners already; a shared score that changed depending on who asked would be a worse lie than none.
- **Raw images are discarded.** After extraction and confirmation we keep the structured transaction record, not the picture of the receipt.

- **Local-first storage.** Data lives in a single-tenant PocketBase instance that the household or operator controls — not a shared third-party processor — and the ledger is append-only and hash-chained, so a correction is recorded rather than overwriting history.
- **PDPA alignment.** We operate consistent with Malaysia's Personal Data Protection Act 2010: users consent to processing, data is used only for the stated purpose, and users can request correction or deletion of their records. We describe this as alignment, not certified compliance, because no audit has been performed.
- **No bank links.** HoneyMoney never asks for banking credentials and holds no open-banking connection, which removes an entire class of risk rather than mitigating it.

5. AI used to build HoneyMoney

We used AI coding assistants (Anthropic's **Claude** and **Claude Code**) during development for scaffolding, refactoring and drafting code across the stack — a Next.js 16 progressive web app, a local-first PocketBase knowledge graph, the Telegram Bot API, and the multi-provider AI layer above. This accelerated delivery, but **every AI-assisted contribution was reviewed, tested and accepted by a human engineer** before it entered the product. Architecture decisions, security-sensitive logic and the financial rules were designed and verified by our team. AI was a tool in the workshop, not the author of record.

6. Limitations and safeguards

- **OCR is not 100% accurate.** Extraction can misread amounts, merchants or dates, particularly on low-quality images or unfamiliar receipt formats. This is exactly why mandatory human confirmation exists: no unverified figure enters a ledger.
- **No automated, binding financial advice.** Honey offers general guidance and behavioural nudges. It does **not** provide licensed financial, investment, tax or legal advice, and it does not execute transactions on a user's behalf. Guidance is framed as suggestion, not instruction.
- **Forecasts are projections, not promises.** The forward-looking warning ("short by the 24th at this pace") is computed from observed spending velocity against configured commitments. It is arithmetic on incomplete information — money spent outside the app is money the projection cannot see.
- **AI can still be wrong or generic.** The numeric allowlist bounds the damage a wrong answer can do; it does not make every answer useful. Users remain in control of every decision about their money.
- **Data minimisation is a standing safeguard.** Limiting what we collect and what we send to a model reduces both privacy risk and the blast radius of any error.

7. Our commitment

If any statement here becomes inaccurate as the product evolves, we will update this disclosure — as we have done in this revision, by removing a capability that had been withdrawn from the product and by naming only the AI providers that are actually wired in. We would rather understate our capabilities and be believed than overstate them and mislead. Candour is not only a competition requirement for us; it is the only way to earn the trust required to help families manage their money.